

To watch the presentation video click [here](#)

Stock Trading Algorithms From Brute Force to Dynamic Programming

By : Loubna Benchakouk

INTRODUCTION

Analysis of 5 different algorithms for stock trading optimization

Tasks covered:

Problem-1

- Brute Force (Task 1) $\rightarrow O(m \times n^2)$
- Greedy (Task 2) $\rightarrow O(m \times n)$
- Single-Transaction DP (Task 3) $\rightarrow O(m \times n)$

Problem-2

- Multi-Transaction DP (Task 5) $\rightarrow O(m \times n \times k)$

Problem-3

- Multi-Transaction DP With Cooldown Period (Task 6) $\rightarrow O(m \times n^2)$

Focus: Algorithm design, performance analysis, and practical trade-offs

Learning objective: Understanding how algorithm choice impacts real-world performance



THE STOCK TRADING PROBLEM

Input: Matrix $A[m \times n]$ where m = stocks, n = days

Goal: Maximize profit from buy/sell transactions

Example:

```
A = [  
  [7, 1, 5, 3, 6],    # stock 1  
  [2, 4, 3, 7, 9],    # stock 2  
  [5, 8, 9, 1, 2],    # stock 3  
  [9, 3, 14, 8, 7]    # stock 4  
]
```

- Stock 4, buy day 2 (\$3) → sell day 3 (\$14) = \$11 profit

Constraints: Must buy before selling, temporal ordering matters

Variations: Single transaction vs. multiple transactions (up to k)

Task 1 - Brute Force

Approach: Check Everything

Algorithm

Three nested loops (stocks → buy days → sell days)

Time Complexity

$O(m \times n^2)$

Example

1000 stocks × 1000 days = ~1 billion operations

Advantages

Simple to understand, guaranteed correct answer

Disadvantages

Extremely slow for large datasets

Best Use

Learning, small problems, algorithm verification

```
def MaxProfitBruteForce(A, m, n):  
    """  
    Brute force algorithm to find the maximum profit from a single buy/sell trans.  
  
    Parameters:  
        A (List[List[int]]): Matrix representing stock prices (m stocks * n days)  
        m (int): Number of stocks (rows)  
        n (int): Number of days (columns)  
  
    Returns:  
        Tuple[int, int, int, int]: (bestStock, bestBuyDay, bestSellDay, maxProfit)  
        All values are 1-based indices.  
        Returns (0, 0, 0, 0) if no profit is possible.  
    """  
  
    # --- edge case: check if the input is empty or missing ---  
    if m <= 0 or n <= 1 or not A or n < 2:  
        return (0, 0, 0, 0)  
  
    # --- edge case: check if matrix has proper dimensions  
    if len(A) != m or any(len(row) != n for row in A):  
        return (0, 0, 0, 0)  
  
    # --- initialize variables to store the best result found ---  
    maxProfit = 0  
    bestStock = 0  
    bestBuyDay = 0  
    bestSellDay = 0  
  
    # --- try every possible stock ---  
    for i in range(m): # stock index (0-based)  
        # --- try every possible buy day ---  
        for j1 in range(n - 1): # buy day  
            # --- try every possible sell day after the buy day ---  
            for j2 in range(j1 + 1, n): # sell day  
                # --- calculate profit for the current transaction ---  
                profit = A[i][j2] - A[i][j1]  
  
            # --- if this transaction gives higher profit, update the result ---  
            if profit > maxProfit:  
                maxProfit = profit  
                # 1-based index  
                bestStock = i + 1  
                bestBuyDay = j1 + 1  
                bestSellDay = j2 + 1  
  
    # --- return result depending on whether any profit was made ---  
    if maxProfit == 0:  
        return (0, 0, 0, 0) # if no profitable transaction found  
    else:  
        return (bestStock, bestBuyDay, bestSellDay, maxProfit)
```

Task 2 - Greedy Optimization

For any sell day, optimal buy day = lowest previous price

Algorithm

Track minimum price while processing days sequentially

Time Complexity

$O(m \times n)$ - 1000× improvement!

Example

1000×1000 dataset = ~1 million operations

Advantages

Fast, elegant, mathematically optimal

Best Use

Production systems, single-transaction scenarios

```
def MaxProfitGreedySolution(A, m, n):  
    """  
    Greedy algorithm to find the maximum profit from a single buy/sell transaction.  
  
    Parameters:  
        A (List[List[int]]): Matrix representing stock prices (m stocks * n days)  
        m (int): Number of stocks  
        n (int): Number of days  
  
    Returns:  
        Tuple[int, int, int, int]: (bestStock, bestBuyDay, bestSellDay, maxProfit)  
        All values use 1-based indexing.  
        Returns (0, 0, 0, 0) if no profit is possible.  
    """  
  
    # --- edge case: check if the input is empty or missing ---  
    if m <= 0 or n <= 1 or not A or n < 2:  
        return (0, 0, 0, 0)  
  
    # --- edge case: check if matrix has proper dimensions  
    if len(A) != m or any(len(row) != n for row in A):  
        return (0, 0, 0, 0)  
  
    # --- initialize variables to store the best result ---  
    maxProfit = 0  
    bestStock = 0  
    bestBuyDay = 0  
    bestSellDay = 0  
  
    # --- iterate over each stock ---  
    for i in range(m):  
        # --- assume the first day's price is the lowest seen so far ---  
        minPrice = A[i][0]      # track minimum price for current stock  
        minDay = 0              # track the day of the minimum price  
  
        # --- iterate through the rest of the days ---  
        for j in range(1, n):  # start from day 1 (second day)  
            profit = A[i][j] - minPrice  # potential profit if sold today  
  
            # --- if this transaction gives better profit, update result ---  
            if profit > maxProfit:  
                maxProfit = profit  
                # 1-based index for stock, buy day, and sell day  
                bestStock = i + 1  
                bestBuyDay = minDay + 1  
                bestSellDay = j + 1  
  
            # --- update minPrice if a new lower price is found ---  
            if A[i][j] < minPrice:  
                minPrice = A[i][j]  
                minDay = j  
  
    # --- return result ---  
    if maxProfit == 0:  
        return (0, 0, 0, 0)  # if no profitable transaction found  
    else:  
        return (bestStock, bestBuyDay, bestSellDay, maxProfit)
```



```
def MaxProfitDynamicProgramming(A, m, n):
    """
    DP approach to find the maximum profit from a single buy/sell transaction.

    Parameters:
        A (List[List[int]]): Matrix representing stock prices (m stocks * n days)
        m (int): Number of stocks
        n (int): Number of days

    Returns:
        Tuple[int, int, int, int]: (bestStock, bestBuyDay, bestSellDay, maxProfit)
        All values use 1-based indexing.
        Returns (0, 0, 0, 0) if no profit is possible.
    """

    # --- edge case: check if the input is empty or missing ---
    if m <= 0 or n <= 1 or not A:
        return (0, 0, 0, 0)

    # --- edge case: check if matrix has proper dimensions
    if len(A) != m or any(len(row) != n for row in A):
        return (0, 0, 0, 0)

    # --- initialize variables to store the best result ---
    maxProfit = 0
    bestStock = 0
    bestBuyDay = 0
    bestSellDay = 0

    # --- loop through each stock ---
    for i in range(m):
        minPrice = A[i][0] # lowest price seen so far for this stock
        minDay = 0 # day when the lowest price occurred

        # --- check each day starting from day 1 ---
        for j in range(1, n):
            currentProfit = A[i][j] - minPrice # profit if we sell today

            # --- update result if we found a better profit ---
            if currentProfit > maxProfit:
                maxProfit = currentProfit
                bestStock = i + 1 # 1-based index
                bestBuyDay = minDay + 1
                bestSellDay = j + 1

            # --- update minPrice if today's price is lower ---
            if A[i][j] < minPrice:
                minPrice = A[i][j]
                minDay = j

    # --- return result ---
    if maxProfit == 0:
        return (0, 0, 0, 0) # if no profitable transaction found
    else:
        return (bestStock, bestBuyDay, bestSellDay, maxProfit)
```

Task 3 - Single Transaction DP

Dynamic Programming Foundation

Approach:

DP principles applied to single-transaction problem

Performance:

$O(m \times n)$ - identical to greedy algorithm

Purpose:

Framework for extensibility, not immediate speed gains

Value:

Foundation for complex multi-transaction scenarios

Comparison:

Same performance as greedy, better structure for extensions

Best Use:

When planning to extend to Tasks 5-6 later

Task 5 - Multiple Transactions DP

Section 1: Algorithm Overview & Input Validation

```
# --- edge case: check if the input is empty or missing ---
if m <= 0 or n <= 1 or not A or k <= 0:
    return []

# --- edge case: check if matrix has proper dimensions
if len(A) != m or any(len(row) != n for row in A):
    return []
```

Section 2: DP Table Initialization

```
# --- initialize DP tables ---
# DP table to store the best profit at each day for up to t transactions
max_profit = []
for transaction_num in range(k + 1):
    day_profits = []
    for day in range(n):
        day_profits.append(0) # start with 0 profit for each day
    max_profit.append(day_profits)

# DP table to remember which (buy, sell) days gave us that profit
best_days = []
for transaction_num in range(k + 1):
    day_pairs = []
    for day in range(n):
        day_pairs.append(None) # no transaction yet
    best_days.append(day_pairs)
```

Section 3: Core DP Algorithm

```
for t in range(1, k + 1):
    max_diff = -prices[0] # best value of (max_profit[t-1][d] - prices[d])
    best_buy_day = 1 # start with day 1 (1-based)

    for day in range(1, n):
        current_profit = prices[day] + max_diff

        # --- option 1: no new transaction today ---
        if max_profit[t][day - 1] >= current_profit:
            max_profit[t][day] = max_profit[t][day - 1]
            best_days[t][day] = best_days[t][day - 1]
        else:
            # --- option 2: sell today using the best past buy day ---
            max_profit[t][day] = current_profit
            best_days[t][day] = (best_buy_day, day + 1) # store 1-based

    # --- update max_diff and corresponding buy day ---
    prev = max_profit[t - 1][day] - prices[day]
    if prev > max_diff:
        max_diff = prev
        best_buy_day = day + 1
```

Task 5 - Multiple Transactions DP

Section 4: Transaction Reconstruction

```
# --- reconstruct transactions for this stock ---
t = k
day = n - 1
while t > 0 and day > 0:
    if best_days[t][day]:
        buy_day, sell_day = best_days[t][day]
        profit = prices[sell_day - 1] - prices[buy_day - 1]
        all_transactions.append((stock_index + 1, buy_day, sell_day, profit))
        day = buy_day - 2 # go to the day before the buy_day
        t -= 1
    else:
        day -= 1
```

Test Result

Selected Non-Overlapping Transactions:

ID	(Stock, Buy, Sell)	Profit
----	--------------------	--------

T1	(2, 1, 2)	7
----	-----------	---

T2	(2, 3, 5)	6
----	-----------	---

T3	(1, 2, 3)	4
----	-----------	---

[(2, 1, 2), (2, 3, 5), (1, 2, 3)] with total profit = 17

Section 5: Global Selection & Return

```
# --- sort all transactions by profit in descending order ---
all_transactions.sort(key=lambda x: x[3], reverse=True)

# --- pick top-k non-overlapping transactions ---
selected = []
for stock, buy1, sell1, profit in all_transactions:
    overlaps = False
    for _, buy2, sell2, _ in selected:
        if not (sell1 <= buy2 or buy1 >= sell2): # overlapping interval
            overlaps = True
            break
    if not overlaps:
        selected.append((stock, buy1, sell1, profit))
        if len(selected) == k:
            break

return selected
```


Task 6 - Multiple Transactions With Cooldown period DP

Problem Overview

Challenge: Cannot buy any stock for c days after selling any stock

Real-World Application: Models trading restrictions, settlement periods, or regulatory constraints

Example: If you sell on day 3 with cooldown=2, you cannot buy until day 6

Algorithm: Dynamic Programming with constraint handling

Section 2: Find all profitable transactions

```
# --- 1: find all profitable transactions ---
# transactions = [(stock_index, buy_day, sell_day, profit, next_day)]
# where next_day = sell_day + cooldown + 1 (1-based index)
transactions = []
for stock_id in range(m):
    prices = A[stock_id]
    for buy_day in range(n - 1):
        for sell_day in range(buy_day + 1, n):
            profit = prices[sell_day] - prices[buy_day]
            if profit > 0:
                buy = buy_day + 1
                sell = sell_day + 1
                next_day = sell + cooldown + 1
                transactions.append((stock_id + 1, buy, sell, profit, next_day))
```

Section 1: Input Validation

Similar to Task-5

Section 3: Dynamic Programming Optimization

```
# --- 3: DP table where dp[day] = (max_profit, best_sequence) ---
dp = [(0, []) for _ in range(n + 2)] # allows access to sell + cooldown + 1
prefix_max = [(0, []) for _ in range(n + 2)]

for t in transactions:
    stock, buy, sell, profit, next_day = t
    next_day = min(next_day, n) # cap to make sure next_day does not exceed n

    # --- compute the best result before or on the buy day ---
    best_profit_before = prefix_max[buy]
    total_profit = best_profit_before[0] + profit
    sequence = best_profit_before[1] + [(stock, buy, sell)]

    if total_profit > dp[next_day][0]:
        dp[next_day] = (total_profit, sequence)

# --- update prefix max at next_day ---
prefix_max[next_day] = max(prefix_max[next_day], dp[next_day], key=lambda x: x[0])
```

Performance Comparison

Complexity Overview

Task	Problem	Algorithm Type	Time Complexity	Problem Description
Task 1	Problem 1	Brute Force	$O(m \times n^2)$	Single transaction
Task 2	Problem 1	Greedy	$O(m \times n)$	Single transaction
Task 3	Problem 1	Dynamic Programming	$O(m \times n)$	Single transaction
Task 5	Problem 2	Dynamic Programming	$O(m \times n \times k)$	Multiple transactions
Task 6	Problem 3	Dynamic Programming	$O(m \times n^2)$	Cooldown constraints

Performance Analysis (1000×1000 Dataset)

Task	Operations Required	Relative Speed	Best Use Case
Task 2/3	~1 million	1000× faster	Production single-transaction
Task 1	~1 billion	Baseline (slowest)	Learning/verification
Task 6	~1 billion	Similar to Task 1	Cooldown scenarios
Task 5 (k=10)	~10 million	100× faster than brute force	Multiple transactions

Algorithm Selection Guide

Brute Force: Learning, datasets $<100 \times 100$, algorithm verification

Greedy: Production systems, single transactions, speed priority

Single DP: Planning future extensions to multi-transaction problems

Multi DP: Realistic trading, portfolio optimization, k-transaction scenarios, cooldown period

Decision Factors:

Dataset size, performance requirements, future extensibility needs

Understanding level, maintenance complexity, correctness guarantees

THANK YOU